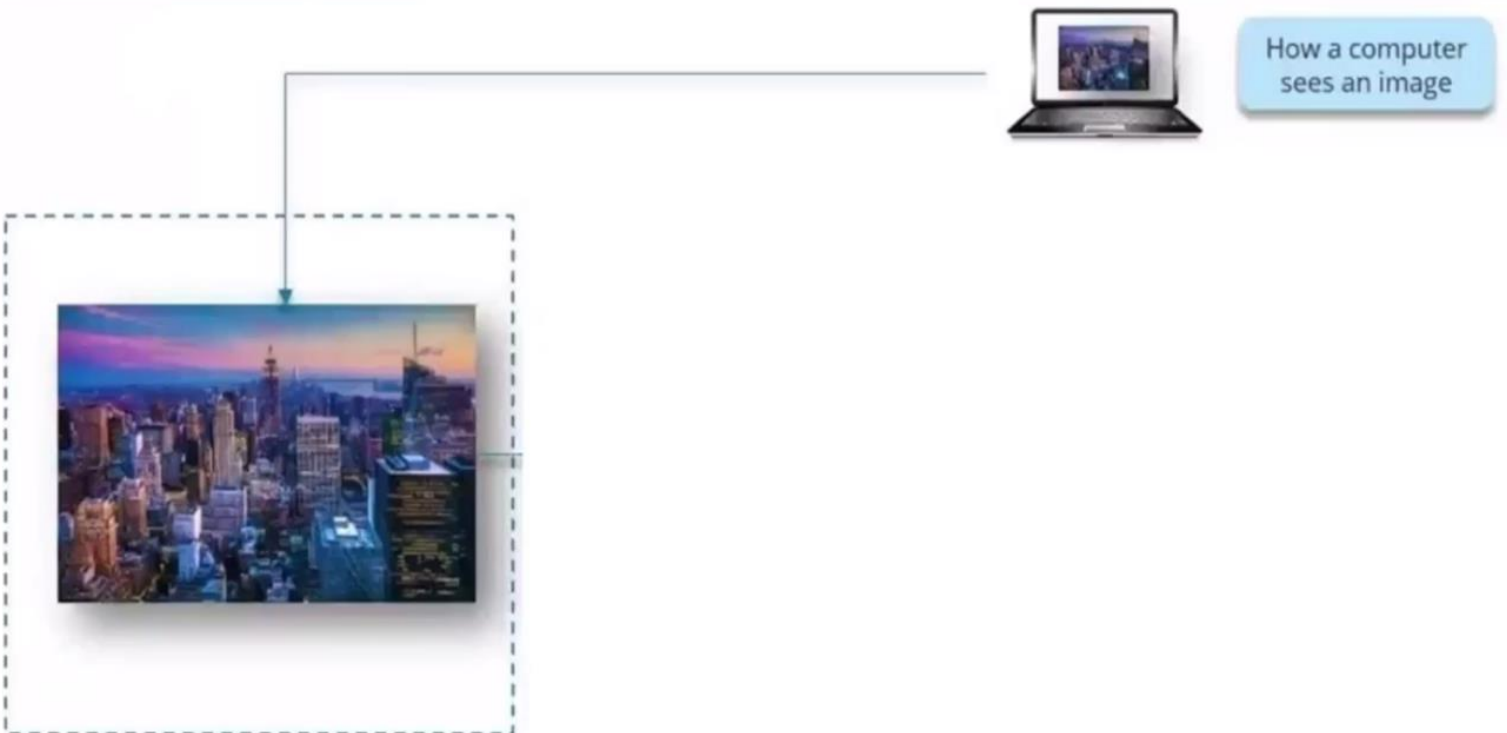


Convolutional Neural Networks

How a Computer Reads an Image



B / W Image 2x2px





2d array



Colored Image 2x2px



3d array



0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Why Not Fully Connected Networks


Let's see why we cannot use fully connected networks for image classifications

The letter "O", zoomed in.


0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

We will start with the image of the letter "O".






Now, because this image is so small, just **6** pixels by **6** pixels...




...into a single column of **36** input nodes...

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0




...and connect the input nodes to a **Hidden Layer**.

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0




So here we have **36** connections from the **36** input nodes to this node in the **Hidden Layer**.

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0



...and each additional node adds an additional **36 Weights** that we need to estimate.

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0



Like I said, because the original image is small (**6x6**) and black and white, something like this could work.

However, if we had a larger image, like **100** pixels by **100** pixels, which is still pretty small compared to real world pictures...

...then we would end up with having to estimate **10,000 Weights per node** in the **Hidden Layer!!!**

So, this method doesn't scale very well.

How CNN Works?

Let's understand how Convolutional Neural Networks Work

Convolutional Neural Networks

STEP 1: Convolution

Step 1 - Convolution Operation
Step 1(b) - ReLU Layer



STEP 2: Max Pooling



STEP 3: Flattening



STEP 4: Full Connection

Step 1 – Convolution

Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter

0	0	1
0	1	0
1	0	0

The first thing a **Convolutional Neural Network** does is apply a **Filter** to the **Input Image**.

Within the filters the values are considered as Weights

Step 1 - Convolution



*

1	0	-1
2	0	-2
1	0	-1



Here also Back propagation will be done to the filter

Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter

3 pixels wide

0	0	1
0	1	0
1	0	0

3 pixels tall

In **Convolutional Neural Networks**, a **filter** is just a smaller square that is commonly **3 pixels by 3 pixels**...

Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter

0	0	1
0	1	0
1	0	0

...and the intensity of each pixel in the filter is determined by **Backpropagation**.

training a **Convolutional Neural Network**, we start with random pixel values...

Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter

1	0	1
0	0	1
1	1	0

...and after training with **Backpropagation**, we end up with something more useful.

Filter

0	0	1
0	1	0
1	0	0

By computing the **Dot Product** between the input and the **Filter**, we can say that the **Filter** is **Convolved** with the input and that's what gives **Convolutional Neural Networks** their name.

Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter

0	0	1
0	1	0
1	0	0

$$\begin{aligned}
 &(0 \times 0) + (0 \times 0) + (1 \times 1) \\
 &+ (0 \times 0) + (1 \times 1) + (0 \times 0) \\
 &+ (1 \times 1) + (0 \times 0) + (0 \times 0)
 \end{aligned}$$

$$= 3$$

+ -2

Feature Map

1		

Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter

0	0	1
0	1	0
1	0	0

...and put the final value into the **Feature Map**.

$$\begin{aligned}
 &(0 \times 0) + (1 \times 0) + (1 \times 1) \\
 &+ (1 \times 0) + (0 \times 1) + (0 \times 0) \\
 &+ (0 \times 1) + (0 \times 0) + (0 \times 0) \\
 &= 1
 \end{aligned}$$

+ -2

Feature Map

1	-1		

Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter


0	0	1
0	1	0
1	0	0

the whole **Feature Map**

+ -2

Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



- you should resize all images to the same dimensions before feeding them into a CNN.
- You can **resize directly** or **use padding** to maintain the aspect ratio.

Size of Feature Map = $N - F + 1$, where N = input image pixels, F = filter pixels

Step 1 (B) – ReLU Layer


ReLU Layer

- ✓ In this layer we remove every negative values from the filtered images and replaces it with zero's
- ✓ This is done to avoid the values from summing up to zero

Rectified Linear Unit (ReLU) transform function only activates a node if the input is above a certain quantity, while the input is below zero, the output is zero, but when the input rises above a certain threshold, it has a linear relationship with the dependent variable

x	$f(x)=x$	F(x)
-3	$f(-3) = 0$	0
-5	$f(-5) = 0$	0
3	$f(3) = 3$	3
5	$f(5) = 5$	5

$$f(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$



Convolutional Neural
Network Tutorial


Feature Map

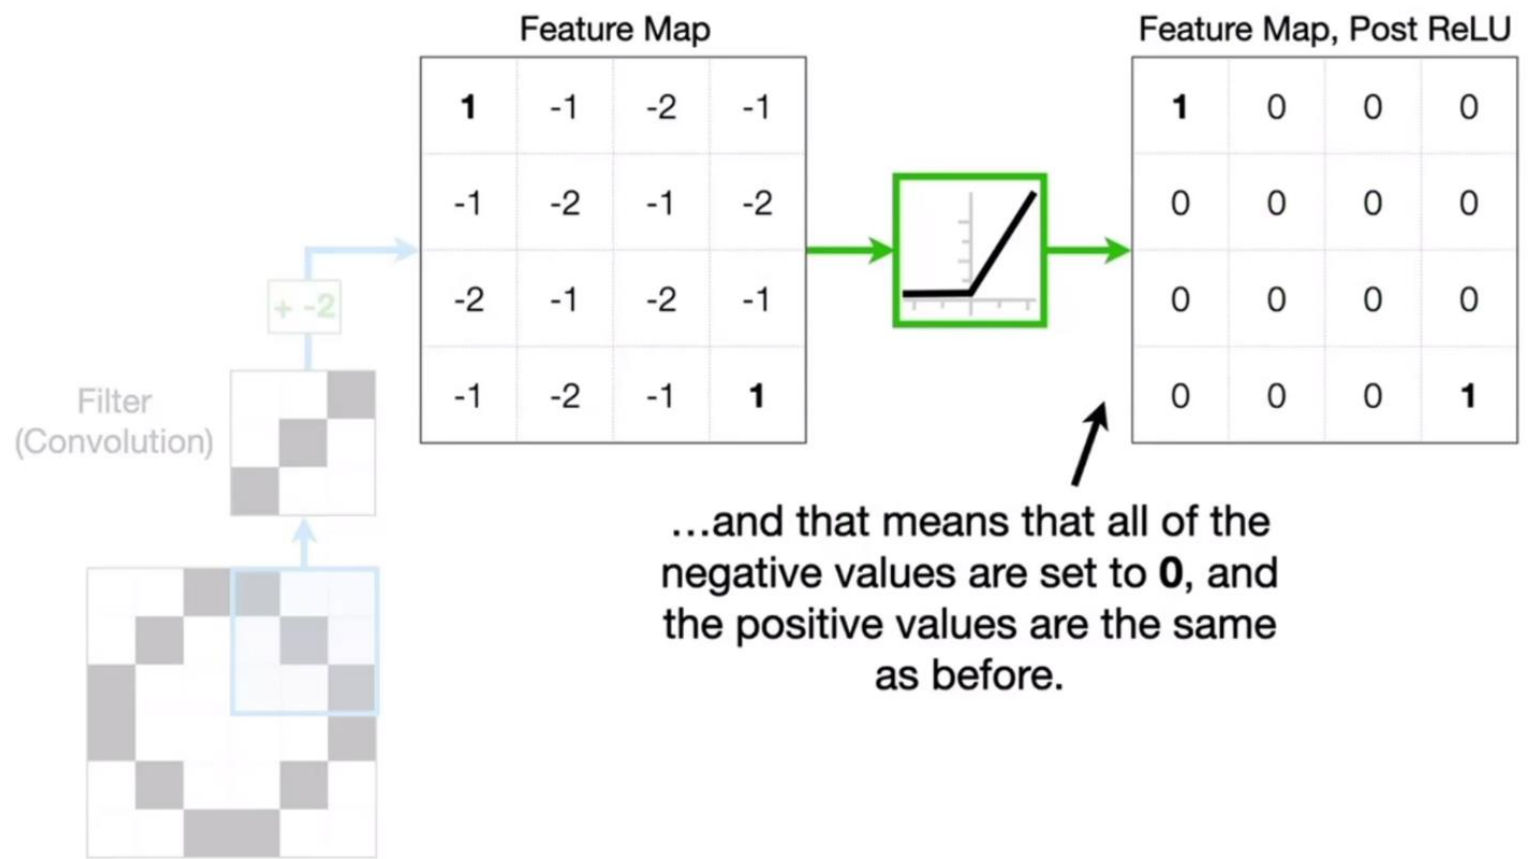
1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



Now, typically, we run the
Feature Map through a **ReLU**
Activation Function...

Filter
(Convolution)






Step 1(B) – ReLU Layer



Step 1(B) – ReLU Layer

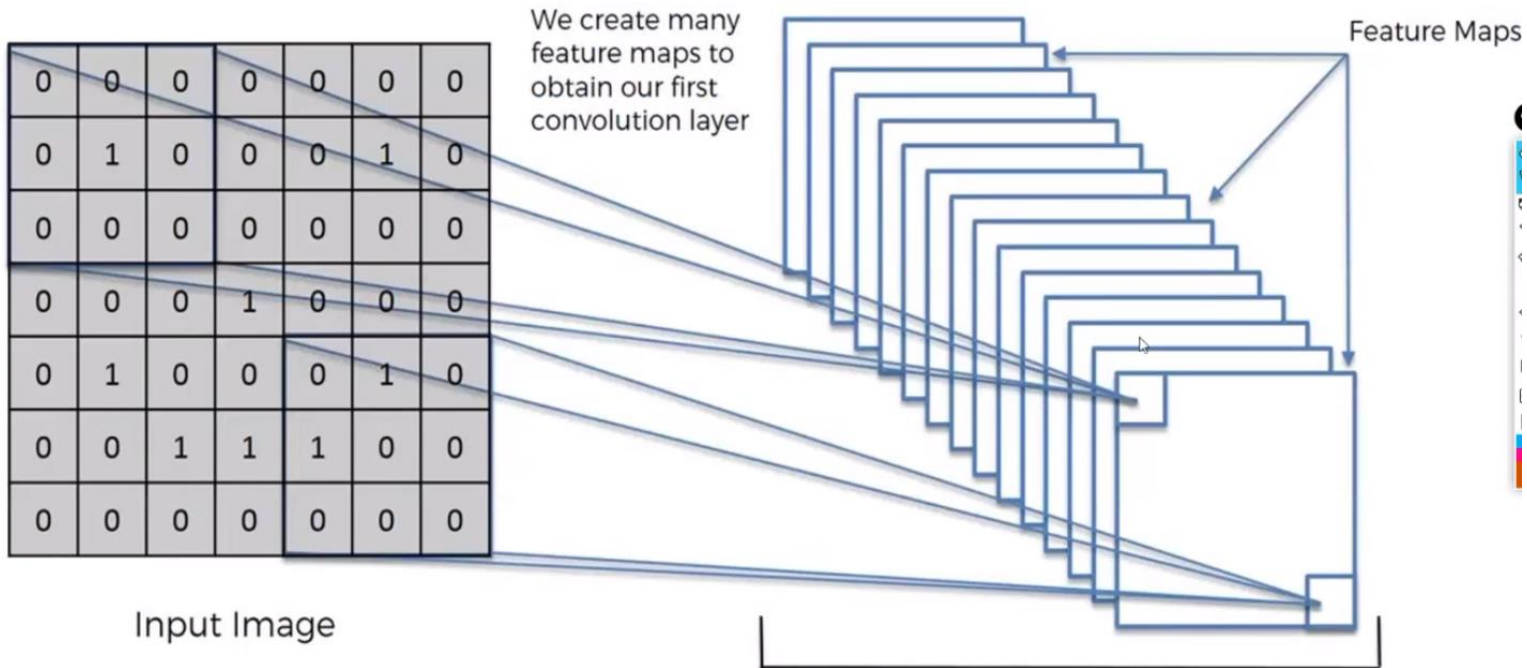


Here ReLU will remove the negitiveness.

Step 1(B) – ReLU Layer



Step 1 - Convolution



Step 2 - Pooling

Pooling Layer

In this layer we shrink the image stack into a smaller size

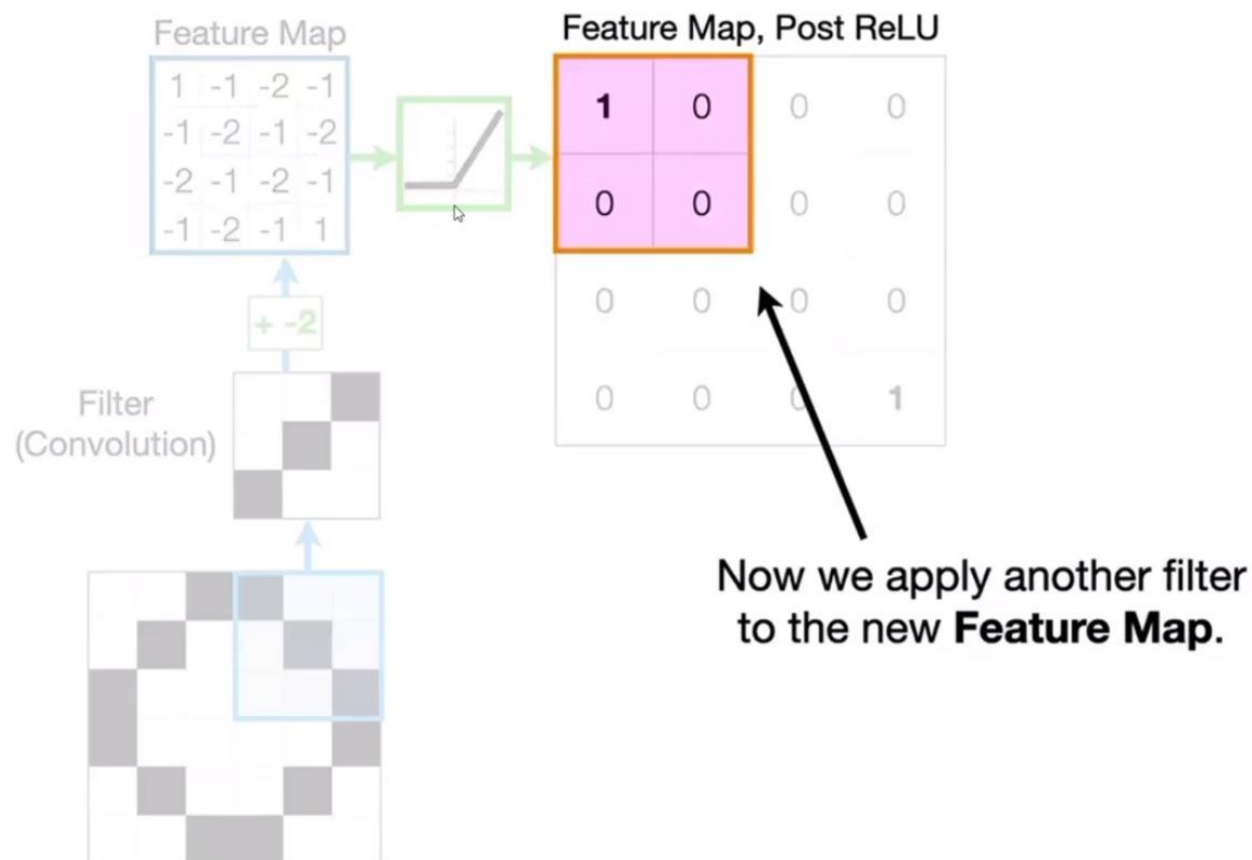
Steps:

1. Pick a window size (usually 2 or 3).
2. Pick a stride (usually 2).
3. Walk your window across your filtered images.
4. From each window, take the maximum value.

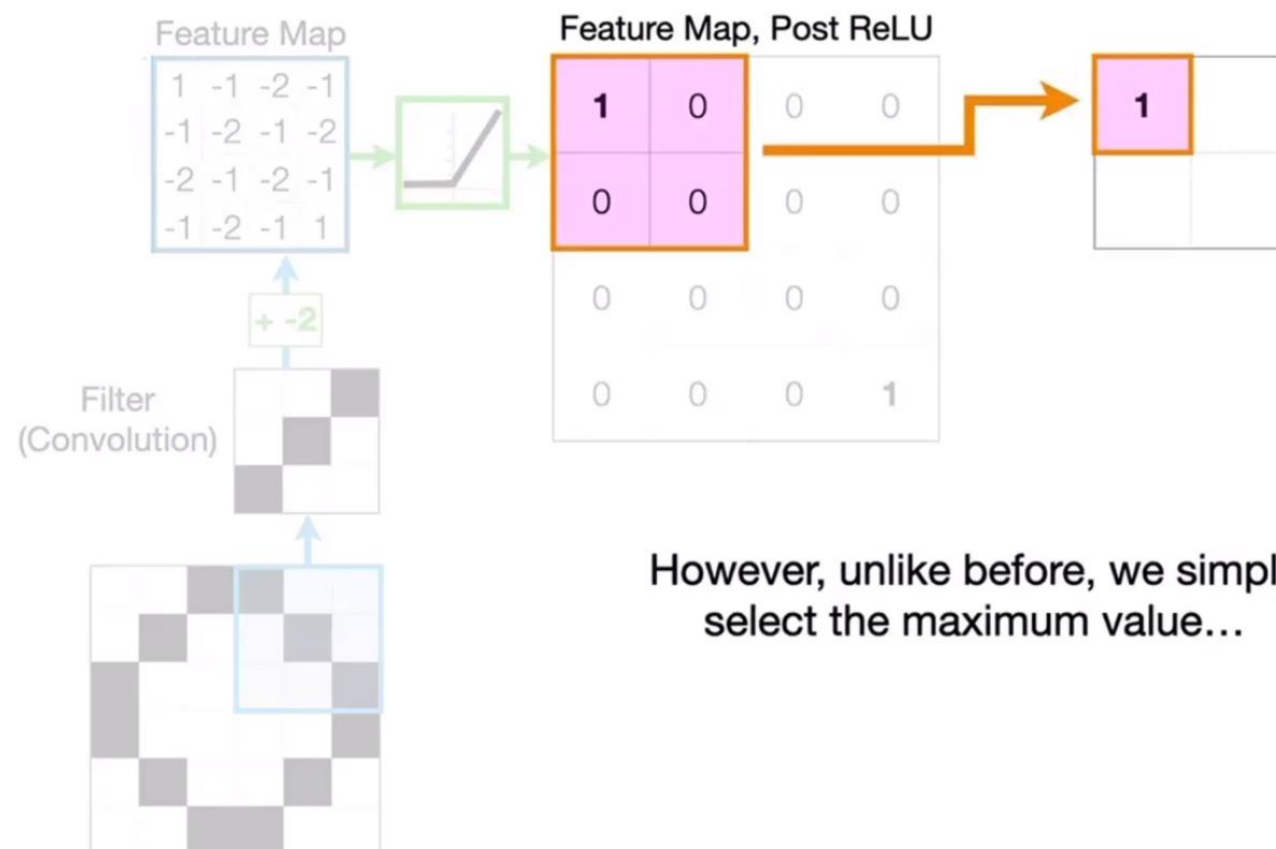
Symbol:



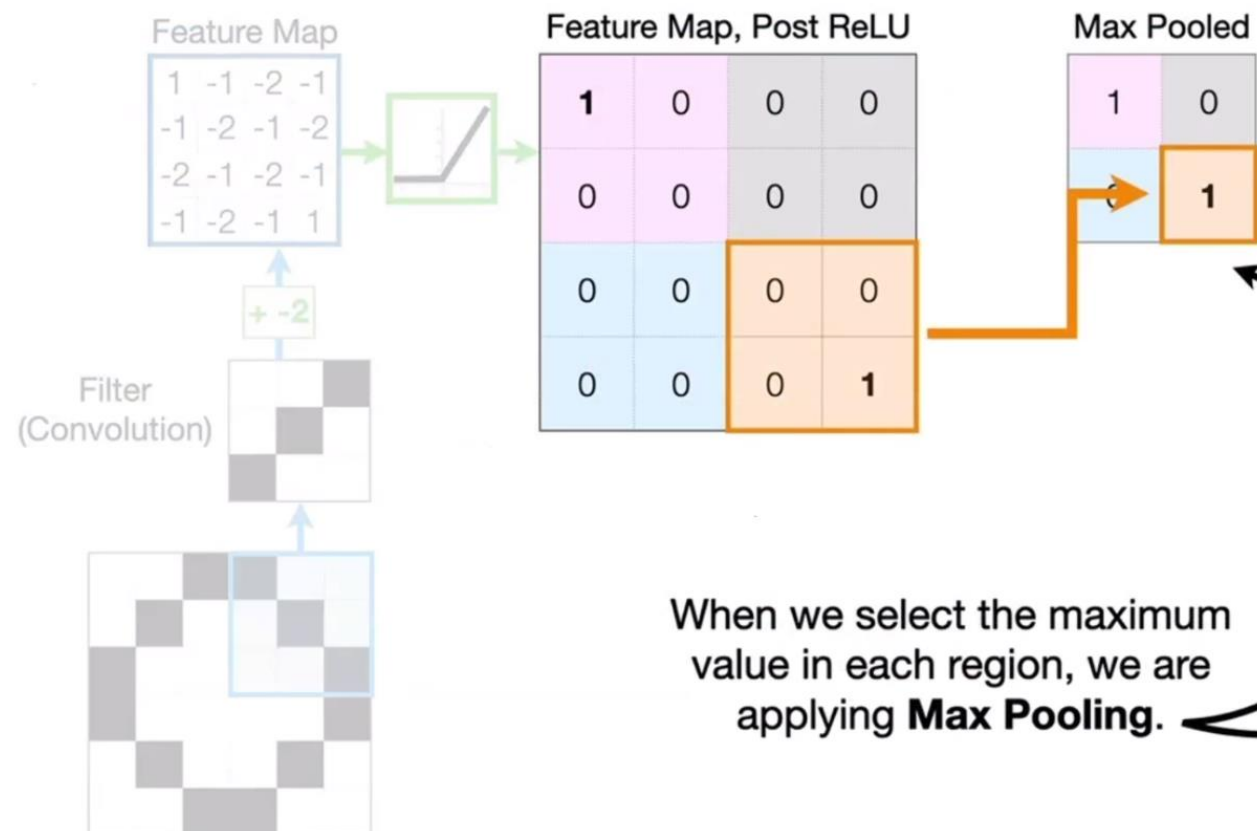
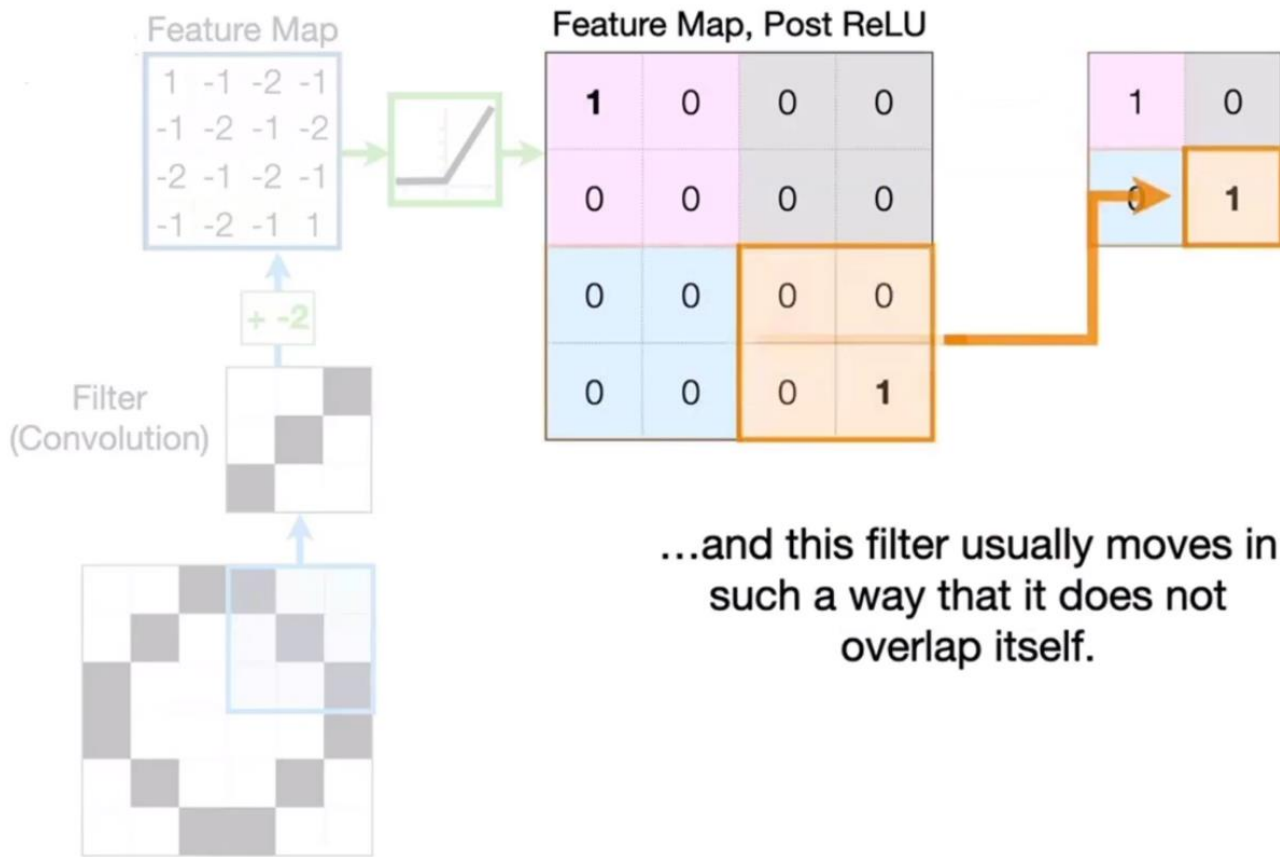
Let's perform pooling with a window size 2 and a stride 2



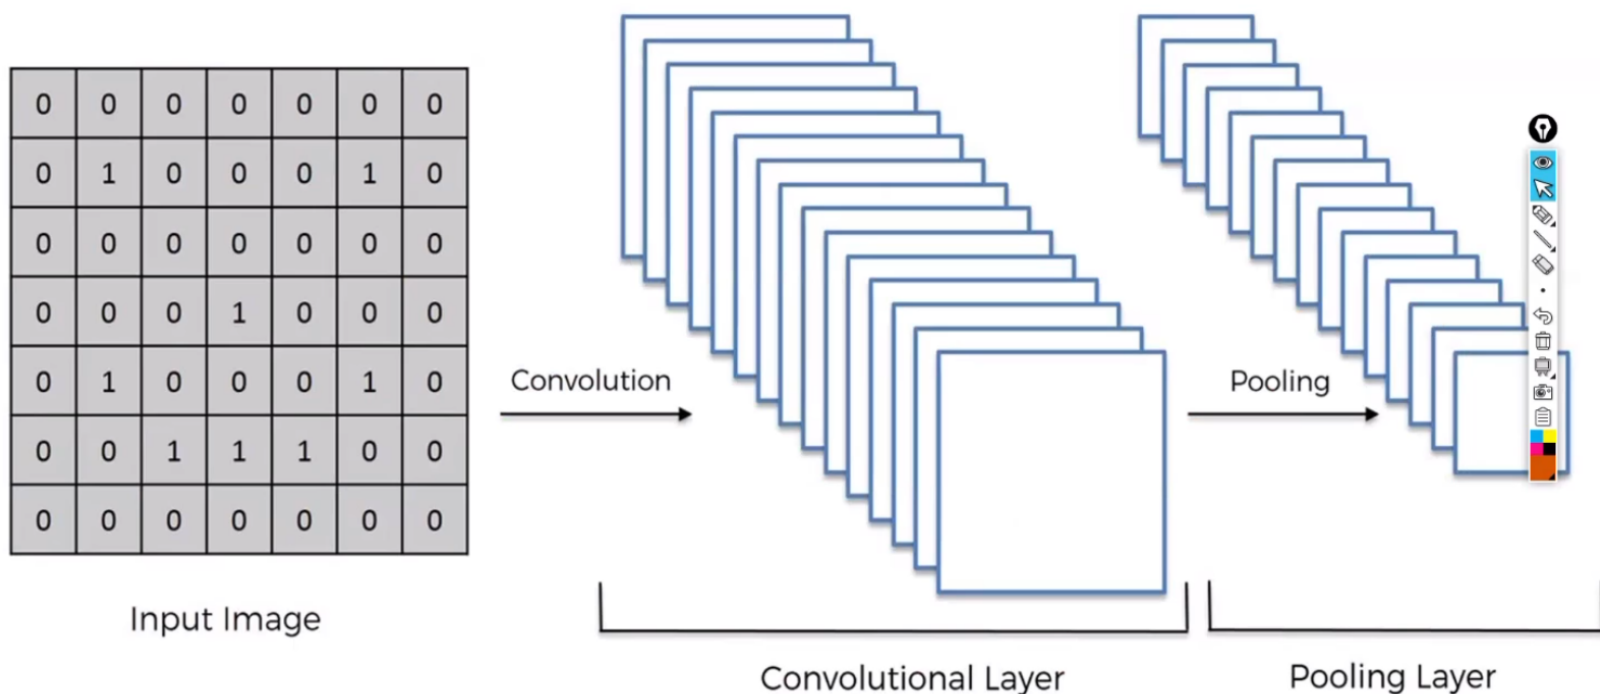
Now we take pool size of 2x2 for which we get max size of 1. This concept is called max pooling. but here it is 2x2 pool size operation is a little different than convolution as it does not overlap.



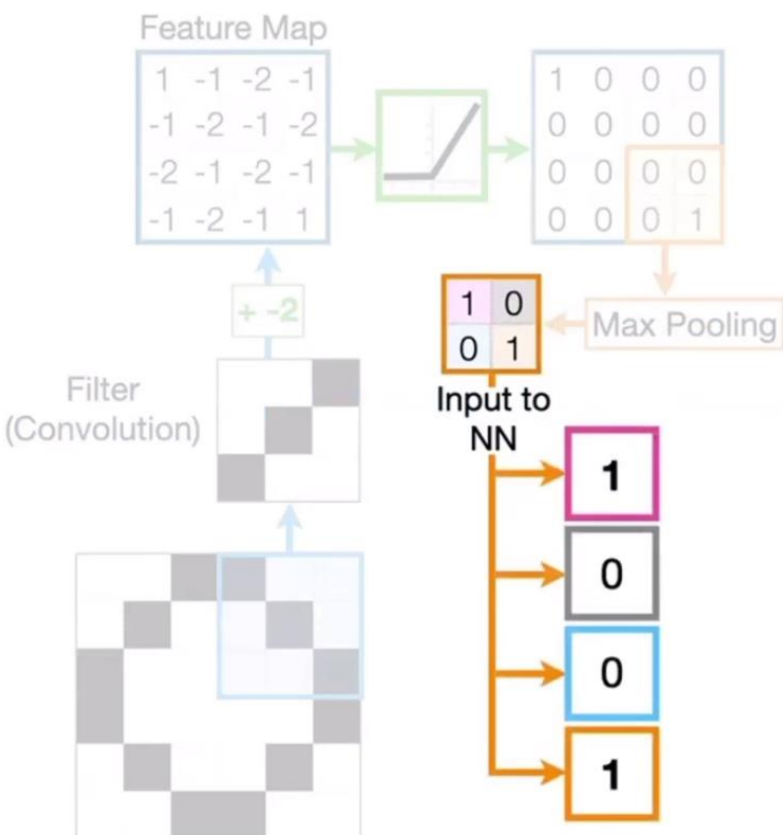
However, unlike before, we simply select the maximum value...



Step 2 - Max Pooling

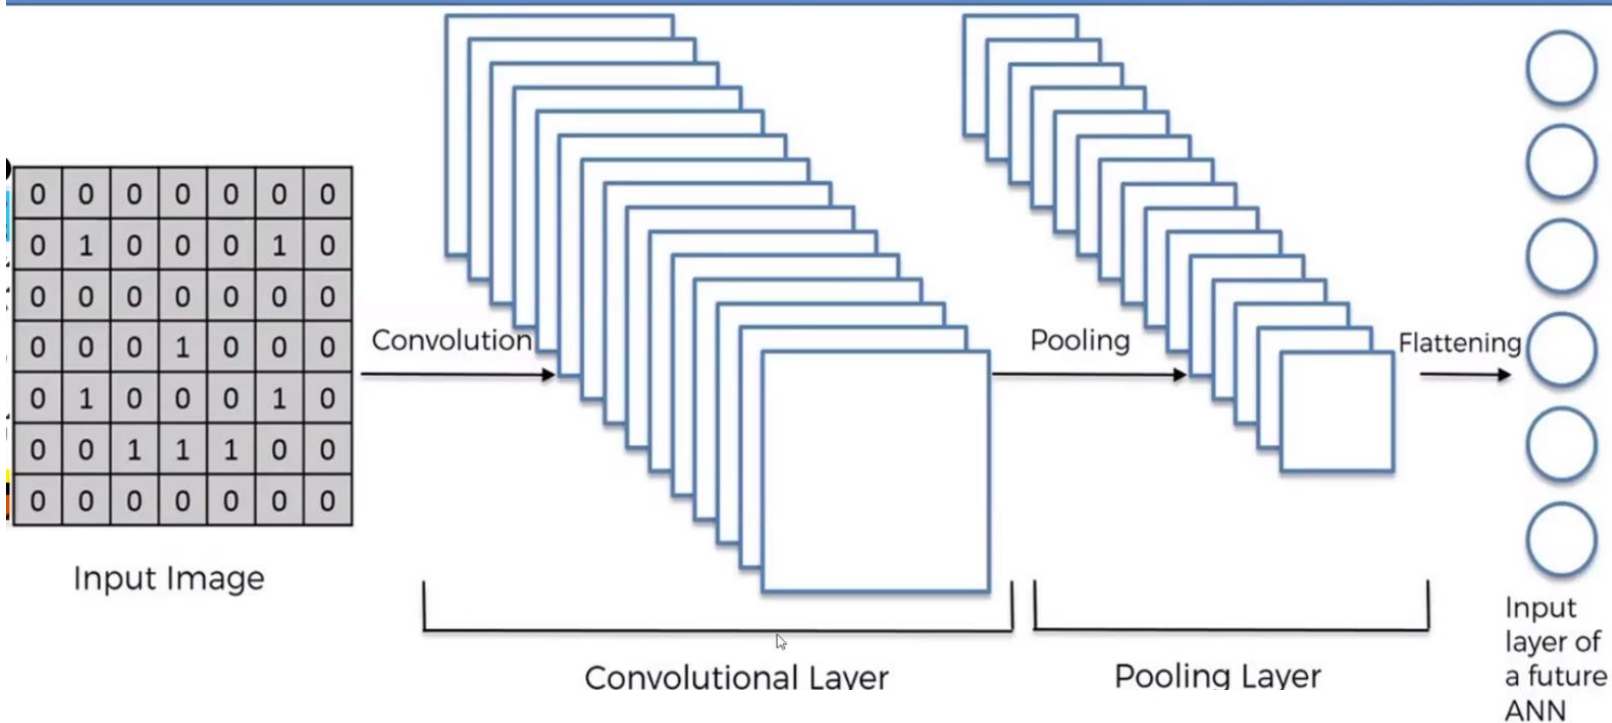


Step 3 - Flattening



Now let's convert the **Pooled Layer** into a column of **Input Nodes**.

Step 3 - Flattening



Step 4 - Full Connection

